



# Assignment\_01

## Assignment 1: Data Collection and Preprocessing for Foundation Model Pre-Training

**Student Name:** [Chen Ting Kuo]

**Date:** February 4, 2026

**Course:** [CSYE7374]

---

### 1. Dataset Sources and Total Size

To ensure domain diversity for foundation model pre-training, we collected data from three distinct high-quality sources: **News (Current Events)**, **Encyclopedic (Factual Knowledge)**, and **Web Text (General Internet Usage)**. The data collection was implemented using the Hugging Face `datasets` library with streaming enabled to handle memory constraints efficiently.

### Data Composition

Domain	Dataset Source	Specific Version / Subset	Samples Collected
News	cc_news	Original (Jan 2017 – Dec 2019)	100,000 documents
Encyclopedic	wikimedia/wikipedia	20231101.en	100,000 documents
General Web	Skylion007/openwebtext	Reddit-sourced WebText	100,000 documents

- **Total Raw Data Volume:** 1.13 GB (approx. 300,000 documents)
  - **Format:** The raw data was aggregated and streamed into a unified `raw_dataset.jsonl` file to ensure persistence and minimize RAM usage during collection.
- 

## 2. Cleaning Strategies and Reasoning

Raw web data is inherently noisy and redundant. We implemented a strict preprocessing pipeline focusing on deduplication and quality filtering.

### 2.1 Deduplication (MD5 Hashing)

Since the `cc_news` dataset contains raw crawls, duplicate articles are common.

- **Strategy:** We utilized **MD5 hashing** to generate a 32-character "fingerprint" for each normalized document.
- **Implementation:** We maintained a `seen_hashes` set in memory. If a document's hash already existed in the set, it was discarded.
- **Reasoning:** Hashing reduces memory footprint significantly and allows for  $O(1)$  lookups, enabling fast exact deduplication.

### 2.2 Normalization and Filtering

- **Normalization:** All text was converted to lowercase, and excessive whitespace (tabs, newlines, multiple spaces) was collapsed into single spaces. This reduces the vocabulary size required and standardizes the input.

- **Length Filtering:** We removed documents with fewer than **50 words**. Short texts often represent navigation menus, error messages, or low-quality snippets that do not provide sufficient context for language modeling.

## 2.3 Cleaning Results

- **Original Documents:** 300,000
- **Duplicates Removed:** 16,358
- **Short Docs Removed:** 9,921
- **Final Cleaned Volume:** **1.10 GB** (273,721 documents)

```
Starting preprocessing pipeline on /Users/zhenting/7374_LLM/raw_dataset.jsonl...
Cleaning: 100% |██████████| 300000/300000 [00:36<00:00, 8127.49it/s]

==== Preprocessing Report ====
Original Docs: 300000
Duplicates Removed: 16358
Short Docs Removed (<50 words): 9921
Final Docs Kept: 273721
Final Dataset Size: 1.10 GB
Cleaned data saved to: /Users/zhenting/7374_LLM/clean_dataset.jsonl
```

## 3. Tokenization Choices

We selected a tokenizer compatible with the **GPT-2** architecture to support autoregressive language modeling tasks.

### 3.1 Tokenizer Configuration

- **Model:** `gpt2` (Hugging Face AutoTokenizer).
- **Algorithm: Byte-Level Byte-Pair Encoding (BPE).** This allows the model to handle out-of-vocabulary words by falling back to byte-level representations, making it robust for diverse web text.
- **Vocabulary Size:** 50,257 tokens.
- **Special Tokens:** Since GPT-2 does not have a native padding token, we mapped `pad_token` to  `eos_token`  (`<|endoftext|>`) to enable compatibility with batch processing.

## 3.2 Chunking and Block Size

- **Block Size: 1024 tokens.**
- **Strategy:** We employed a **Concatenation and Slicing** strategy.
  1. All documents were tokenized and appended with an `EOS` token.
  2. The tokens were concatenated into a continuous stream.
  3. The stream was sliced into fixed-size blocks of 1024 tokens.
  4. The final incomplete block was discarded to ensure consistent tensor shapes.
- **Reasoning:** This approach maximizes training efficiency by eliminating padding inside training blocks and allowing the model to learn context across document boundaries.

## 3.3 Storage (Sharding)

To prevent memory overflows during the conversion of 1GB+ text into integer tensors, we implemented **Sharding**:

- Processed tensors were saved into multiple `.pt` files (`shard_0.pt`, `shard_1.pt`, ...), each containing approximately 50,000 blocks.
- **Final Token Count:** ~266 Million tokens.

```
Loading tokenizer: gpt2...
Tokenizing with gpt2 (Block size: 1024)...

Tokenizing:  0% | 0/273721 [00:00<?, ?it/s] Token indices specified maximum sequence length for this model (1123 > 1024). Running this sequence through the model will result in truncation.
Tokenizing: 29% | 79668/273721 [00:44<13:21, 242.13it/s]

50004
tensor([ 8117,   338,   257,   ..., 16667, 28841,   365])
tensor([[ 8117,   338,   257,   ...,    13,  2008, 12537],
       [ 299,    88, 21101,   ...,    85,  4763, 47735],
       [21421,    13,   299,   ...,   717,  5545,  3652],
       ...,
       [40138,   410, 40138,   ..., 45630,   272, 14549],
       [19322,   784, 44873,   ..., 3104,   784,   556],
       [ 2516,   286,   277,   ..., 16667, 28841,   365]])
Saved shard 0 to /Users/zhenting/7374_LLM/tokenized_data/shard_0.pt: shape torch.Size([50004, 1024])

Tokenizing: 45% | 122192/273721 [01:34<02:39, 949.94it/s]
```

## 4. Data Loader Implementation Details

We implemented a custom PyTorch `IterableDataset` (`GPTDataset`) to handle the large-scale data efficiently without loading the entire dataset into RAM.

### 4.1 Streaming Architecture

The data loader does not preload data. Instead, it iterates through the `.pt` shard files one by one. Only one shard (approx. 50MB–100MB) resides in memory at any given time.

### 4.2 Double Shuffling Strategy

1. **Shard-Level Shuffle:** The order of `.pt` files is randomized at the start of each epoch.
2. **Sample-Level Shuffle:** Once a shard is loaded into memory, its internal indices (rows) are shuffled before yielding.

```
Loading shard: shard_1.pt
Batch 0 shape: torch.Size([8, 1024])
Batch 1 shape: torch.Size([8, 1024])
Batch 2 shape: torch.Size([8, 1024])
Data Loader test passed!
```

---

## 5. Challenges Encountered

### Memory Management (RAM)

- **Issue:** Initially, appending all 300,000 raw text strings to a Python list caused the kernel to crash due to high memory overhead.
- **Solution:** We switched to a **streaming approach**, writing each document directly to a `.jsonl` file immediately after fetching (`stream_to_file`). Similarly, during tokenization, we periodically flushed processed tokens to disk (`shard_X.pt`) to clear the buffer.

---

## 6. Reflections on Preprocessing Impact

This assignment highlighted the critical trade-offs between **data quality** and **computational cost**.

1. **Impact of Deduplication:** While MD5 deduplication added a processing overhead, removing ~16,000 duplicates ensures the model doesn't memorize repeated text, which promotes better generalization.
2. **Normalization Trade-off:** Aggressive normalization (lowercasing) simplifies the vocabulary but loses semantic nuances (e.g., "Apple" the company vs. "apple" the fruit). For a larger-scale model, a cased tokenizer might be preferable.
3. **Efficiency of Binary Formats:** Storing data as `.pt` tensors rather than raw text significantly accelerated the Data Loader. The CPU no longer needs to tokenize text on-the-fly during training, preventing the data loading step from becoming a bottleneck for the GPU.